

# Homework 1 - System Calls and Virtual Memory

Due: October 15th, 2018, 11:00pm

GitHub page: <https://github.com/SudarsunKannan/CS519/blob/master/homework1/homework1.md>

In the second homework, we will explore the cost of system calls and the page fault cost. The homework will consist of three parts: part 1 and part 2 will require you to develop a benchmark to measure the cost of a system call, whereas the part 3 will involve reducing the cost of page fault handling by modifying the OS virtual memory management.

Linux 4.17 kernel source code can be found here

```
https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/?h=linux-4.17.y
```

## PART 1

The first step would be to write a simple benchmark to measure the cost of the system call; you will add a new dummy system call (hello\_kernel) to your OS in the Linux kernel's [mm/mmap.c](#) file. A couple of references below show how to add a new system call; there are thousands of other references online.

[Reference 1](#) [Reference 2](#)

Once you have added a system call, measure the cost of invoking a system call from userspace by invoking the system call few hundred times and measuring the average latency.

NOTE: Make sure that your system call is getting called by adding a **printk()** within your dummy system call. You should remove the print statement after confirming that the new system is getting invoked inside the kernel.

## PART 2

In the second part, we will study the cost of page fault. Remember that Linux OS allocates pages **on demand** (the first time a page is accessed (or touched) after allocation). First, you will allocate a large 2GB memory-mapped region which is also page-aligned. Next, you will touch (write) the first byte of each page sequentially and measure the average cost of accessing a page (i.e., page fault cost).

**Note: For allocating memory-mapped region**

Use `mmap()` system call that allocates page-aligned memory. When using `mmap()`, you should map an anonymous memory and not a file backed memory. See the following link for more details how to allocate anonymous memory.

[Reference 1](#) [Reference 2](#)

## PART 3

The part 3 of homework is an attempt to reduce the cost of page fault handling cost. Every time a page fault occurs, instead of allocating only one page, you are required to allocate two or more pages after a page fault. For example, one page for the actual faulting address (say, addr X) and the next page (say, addr X + 4096). For this, you will modify the OS virtual memory fault handler.

The page faults for anonymous memory are handled in the following function inside the [mm/memory.c](#) source file.

```
int do_anonymous_page(struct vm_fault *vmf)
```

This function first checks whether the page fault is read or write fault (due to read or write access), then (1) allocates a page, (2) creates a new page table entry (PTE), and (3) adds the PTE to the page table.

We will walk through this function in the class.

## Resources

In this class, we will use the QEMU-based virtual machine to test the modified OS. QEMU VMs run either on a bare-metal OS or even inside another virtual machine. More details on QEMU can be obtained [here] (<https://www.qemu.org/>).

For students new to QEMU or hacking kernel, we have created a set of instructions about how to compile a custom kernel (OS) and how to run the OS using QEMU. Detailed step-by-step instructions can be found [here](#).

```
https://github.com/SudarsunKannan/CS519
```

## Computing Resource

If you need access to a development environment and cannot use your laptop, please send me an email ([sudarsun.kannan@cs.rutgers.edu](mailto:sudarsun.kannan@cs.rutgers.edu)).

## Starting Early

This is a significant-but-essential homework for understanding the basics of OS virtual memory. Please start working on this homework early. If you have questions, make sure to ask them during office hours.